

Machine Learning for Engineering and Science Applications



Machine Representation of Numbers
Overflow
Underflow

Machine Arithmetic

- The derivations and expressions in the previous slides assume real number arithmetic.
 - That is, all calculations are assumed to take place perfectly
- However, in practice, the machine only stores numbers to a finite precision
- This arithmetic is called “Finite Precision Arithmetic”, “Machine Arithmetic”
 - Or “Floating Point arithmetic” in the special case of floating point numbers
- This can have surprisingly important consequences

The Ariane 5 disaster

- Ariane 5 is a European launch vehicle.
- Its very first test was on June 4th, 1996.
- 37 seconds after launch the rocket turned by 90 degrees incorrectly.
- The boosters were ripped apart and the vehicle self destructed.
- The loss was about \$500 million
- The reason was found to be software failure primarily due to forgetting to account for finite precision arithmetic.



How a machine stores numbers

- Machines have a **finite** number of “bits”
 - Think of them as boxes where numbers are stored in binary (0 or 1)
- Computers usually use the binary (Base-2) system

- Integer representation :

$$(10101101)_2 = 2^0 + 2^2 + 2^3 + 2^5 + 2^7 = (173)_{10}$$




- Signed-integer representation : Use a sign bit (1 for negative)
- Example: With 16-bits, $-(173)$ will be represented as



This means there is a max and min number that can be represented on the computer

Floating point representation

- Real numbers can also be represented in binary
 - e.g $5.5 = 4 + 0 + 1 + 0.5 = (1 \times 2^2) + (0 \times 2^1) + (1 \times 2^0) + (1 \times 2^{-1}) = (101.1)_2$
- The **floating-point format** is to use the scientific notation $\pm s \times b^e$
 - S is the significand, b the base we are using, and e the exponent
 - Eg : $0.001234 = 1.234 \times 10^{-3}$
- Binary floating point numbers are represented as $\pm (1 + f) \times 2^e$
 - Ex : $(5.5)_{10} = (101.1)_2 = 1.011 \times 2^{-2} = (1 + 0.011) \times 2^{-2}$
 - f is the mantissa and e the exponent
 - Each of f and e require separate bins to store
- For a 64-bit storage scheme



If you have understood that there is a min and max number on the computer you can skip this slide

Double precision – Range

- IEEE Double precision – 64 bits (8 bytes) are used for storing floating point numbers
 - Is the standard amount of precision used for much of scientific computation
 - MATLAB uses this by default
- Since there are only limited “boxes” for storing the exponent, there is a max/min positive number that can be represented
 - 11 bits for the signed exponent \Rightarrow Range of the exponent is from -1022 to 1023
 - Largest number = $1.111 \dots 111 \times 2^{+1023} \approx 1.8 \times 10^{308}$
 - Smallest number = $1.000 \dots 000 \times 2^{-1022} \approx 2.2 \times 10^{-308}$

```
>> format long
>> realmax

ans =

    1.797693134862316e+308

>> realmin

ans =

    2.225073858507201e-308
```

Double precision – Precision

Consider the following example

```
>> sqrt(2)
```

```
ans =
```

```
1.414213562373095
```

```
>> err = 1.e-14
```

```
err =
```

```
1.0000000000000000e-14
```

```
>> sqrt(2)+err
```

```
ans =
```

```
1.414213562373105
```

```
>> err = 1.e-16
```

```
err =
```

```
1.0000000000000000e-16
```

```
>> sqrt(2)+err
```

```
ans =
```

```
1.414213562373095
```

```
>> a = 1; b = -1; err = 1.e-16;
```

```
>> a+b+err
```

```
ans =
```

```
1.0000000000000000e-16
```

```
>> a+err+b
```

```
ans =
```

```
0
```

- Happen because even mantissa is limited to 52 bits
- Double Precision (called Machine epsilon) is given by $2^{-52} \approx 2.22 \times 10^{-16}$

```
>> eps
```

```
ans =
```

```
2.220446049250313e-16
```

Underflow and overflow

- **Underflow** happens due to numbers near zero being “rounded off” to zero

```
>> err = 1.e-16  
  
err =  
  
1.0000000000000000e-16  
  
>> sqrt(2) + err  
  
ans =  
  
1.414213562373095
```

Overflow

Underflow

Overflow

- **Underflow** might result in divide by zero errors even when the denominator is finite
- **Overflow** happens when the number being computed overshoots the max possible

■ Consider $\frac{\exp(x_1)}{\exp(x_1) + \exp(x_2)}$

```
>> x = [5000 5000];  
>> soft = exp(x(1))/(exp(x(1)) + exp(x(2)))  
  
soft =  
  
NaN
```


```
>> z = x - max(x);  
>> soft = exp(z(1))/(exp(z(1)) + exp(z(2)))  
  
soft =  
  
0.5000000000000000
```

- Can be reformulated to account for overflow

The Ariane 5 disaster was due to an overflow problem

The Ariane 5 disaster

- The internal Inquiry board reported that this was due to software error in the Inertial Reference System (**SRI** -- Système de Référence Inertielle)
- The SRI used a floating point variable BH to determine the orientation of the rocket.
- This variable was stored as a 64-bit floating point
 - It had to be converted to a 16-bit signed integer
- This caused no problems for the previous version (Ariane 5) as the values were below the max
- After 37s of flight in Ariane 5, these values were exceeded
 - Accelerations were higher in Ariane 5 than Ariane 4
- In the words of the report –
 - The internal SRI software exception was caused during execution of a data conversion from 64-bit floating point to 16-bit signed integer value. The floating point number which was converted had a value greater than what could be represented by a 16-bit signed integer. This resulted in an Operand Error.



Condition number

- Limited precision has many unexpected results

```
s=0;  
for i = 1:10000  
    s = s+0.0001;  
end  
disp(s)
```

```
>> PrecisionTest  
0.999999999999906
```

Precision effects propagate

- This can lead to problems when systems of linear equations are solved

A =

```
1.0000000000000000 2.0000000000000000  
2.0000000000000000 4.000000000100000
```

x =

```
1  
-1
```

```
>> b = A*x
```

```
b =  
-1.0000000000000000  
-2.000000000100000
```

```
>> inv(A)*b
```

```
ans =  
0.999999999950000  
-0.999999999975000
```

- Some matrices are close to singular

- Small errors can be magnified by them

- This is measured by condition number

- In general $\text{cond}(A) = \|A\| \|A^{-1}\|$

- For symmetric matrices, condition number is the ratio of the largest to smallest eigenvalue

- $\text{cond}(A) = \max_{i,j} \frac{|\lambda_i|}{|\lambda_j|}$. Higher condition number means poorly conditioned

```
>> b1
```

```
b1 =  
-1.0100000000000000  
-2.0100000000000000
```

```
>> inv(A)*b1
```

```
ans =  
1.0e+08 *  
-1.999999844569314  
0.999999917234647
```

```
>> cond(A)
```

```
ans =  
2.499969680591895e+11
```